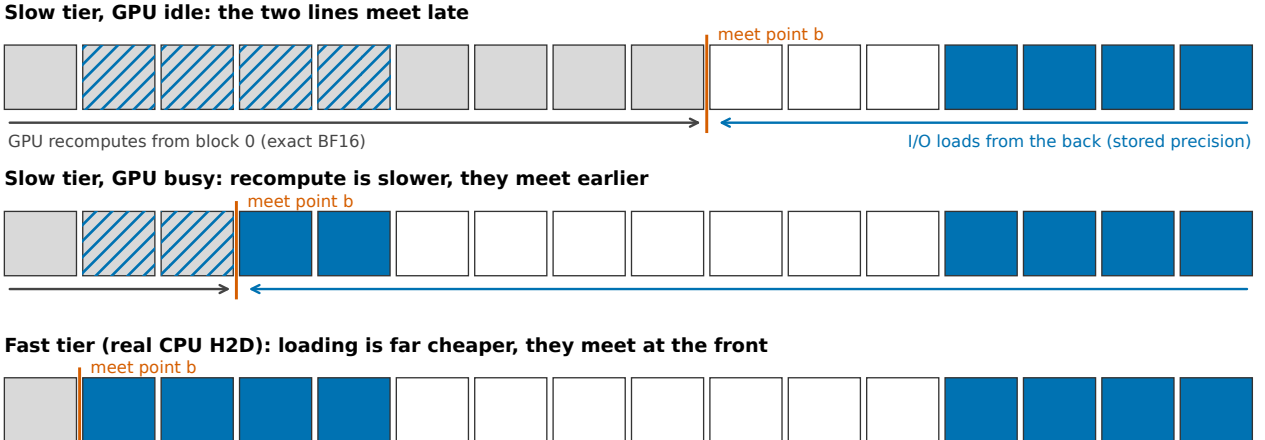
Slow tier, GPU idle: the two lines meet late
meet point b
GPU recomputes from block 0 (exact BF16)
I/O loads from the back (stored precision)
Slow tier, GPU busy: recompute is slower, they meet earlier
meet point b
Fast tier (real CPU H2D): loading is far cheaper, they meet at the front
meet point b
recomputed (exact BF16)
loaded BF16
loaded low precision (read)
stored low precision, never read