1. K0 (FP8): does all-FP8 restore faster?
E1 wins 7/8 cells (Llama), 6/8 (Qwen) -> K0 false, continue

2. Step 1 (FP8 endpoints): is uniform FP8 enough?
Llama K1 = yes; Qwen K1 = no, K1s = no

3. Step 2b (Qwen): does an 8-bit codec change fix it?
E1-K0 pass = no, E1-INT8 pass = yes -> K3 = yes: FP8 tier needs no placement

4. K0 (INT4) + Step 1' (INT4 endpoints): is uniform INT4 enough?
K0: E1 wins 6/8 (Llama), 6/8 (Qwen) -> K0 false. Quality eps, all 4 readings: Llama yes, Qwen yes.
Speed-not-worse fails only on the real CPU tier (Llama cpu/busy; Qwen cpu/busy, cpu/idle) -> K1 = K1s = no

5. Route (03 Sec.1.1): uniform INT4 'not enough' for both -> Step 2 on Llama
Qwen INT4 mixed configs split out (not part of the decision)

6. Step 2 (INT4, Llama): F25/M25/B25 vs R25
stable gap (cond. 1-2) = no -> K2 = yes: position adds nothing; STOP

not reached: FP8 Step 2, Step 3 (sensitivity-aware rule S25)