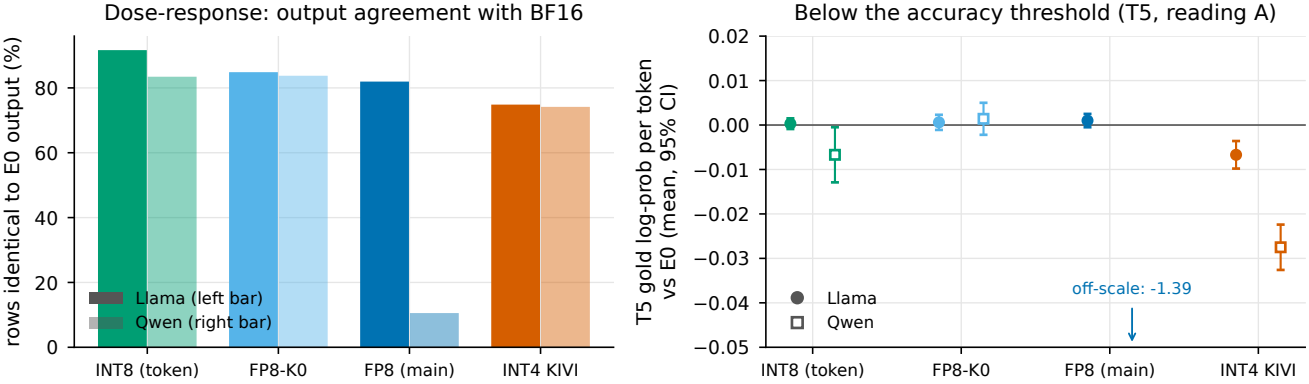
Dose-response: output agreement with BF16
rows identical to E0 output (%)
0
20
40
60
80
Llama (left bar)
Qwen (right bar)
INT8 (token)
FP8-K0
FP8 (main)
INT4 KIVI
Below the accuracy threshold (T5, reading A)
T5 gold log-prob per token
vs E0 (mean, 95% CI)
0.02
0.01
0.00
−0.01
−0.02
−0.03
−0.04
−0.05
Llama
Qwen
off-scale: -1.39
INT8 (token)
FP8-K0
FP8 (main)
INT4 KIVI