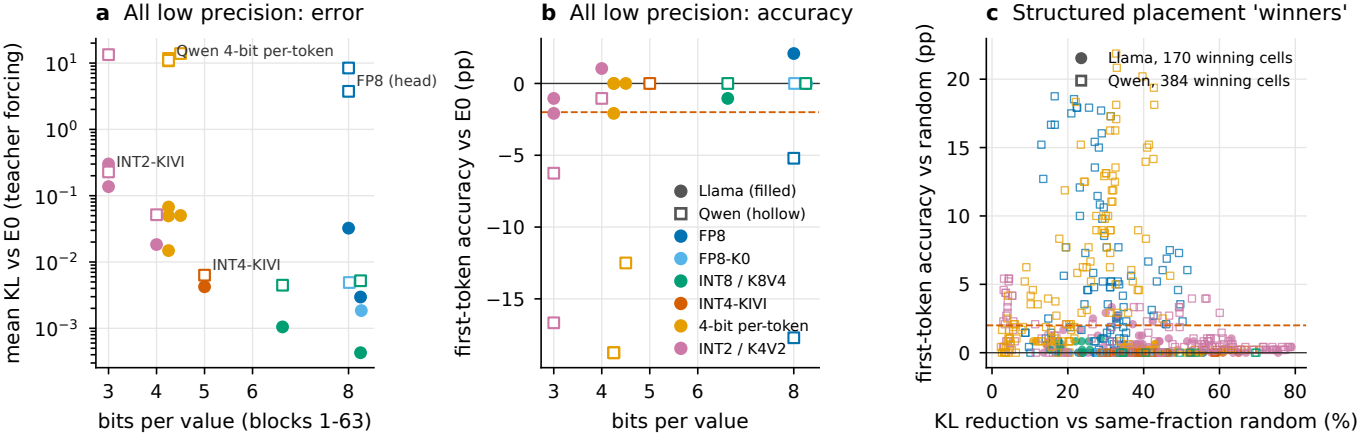
a All low precision: error
mean KL vs E0 (teacher forcing)
bits per value (blocks 1-63)
Qwen 4-bit per-token
FP8 (head)
INT2-KIVI
INT4-KIVI
b All low precision: accuracy
first-token accuracy vs E0 (pp)
bits per value
Llama (filled)
Qwen (hollow)
FP8
FP8-K0
INT8 / K8V4
INT4-KIVI
4-bit per-token
INT2 / K4V2
c Structured placement 'winners'
first-token accuracy vs random (pp)
KL reduction vs same-fraction random (%)
Llama, 170 winning cells
Qwen, 384 winning cells