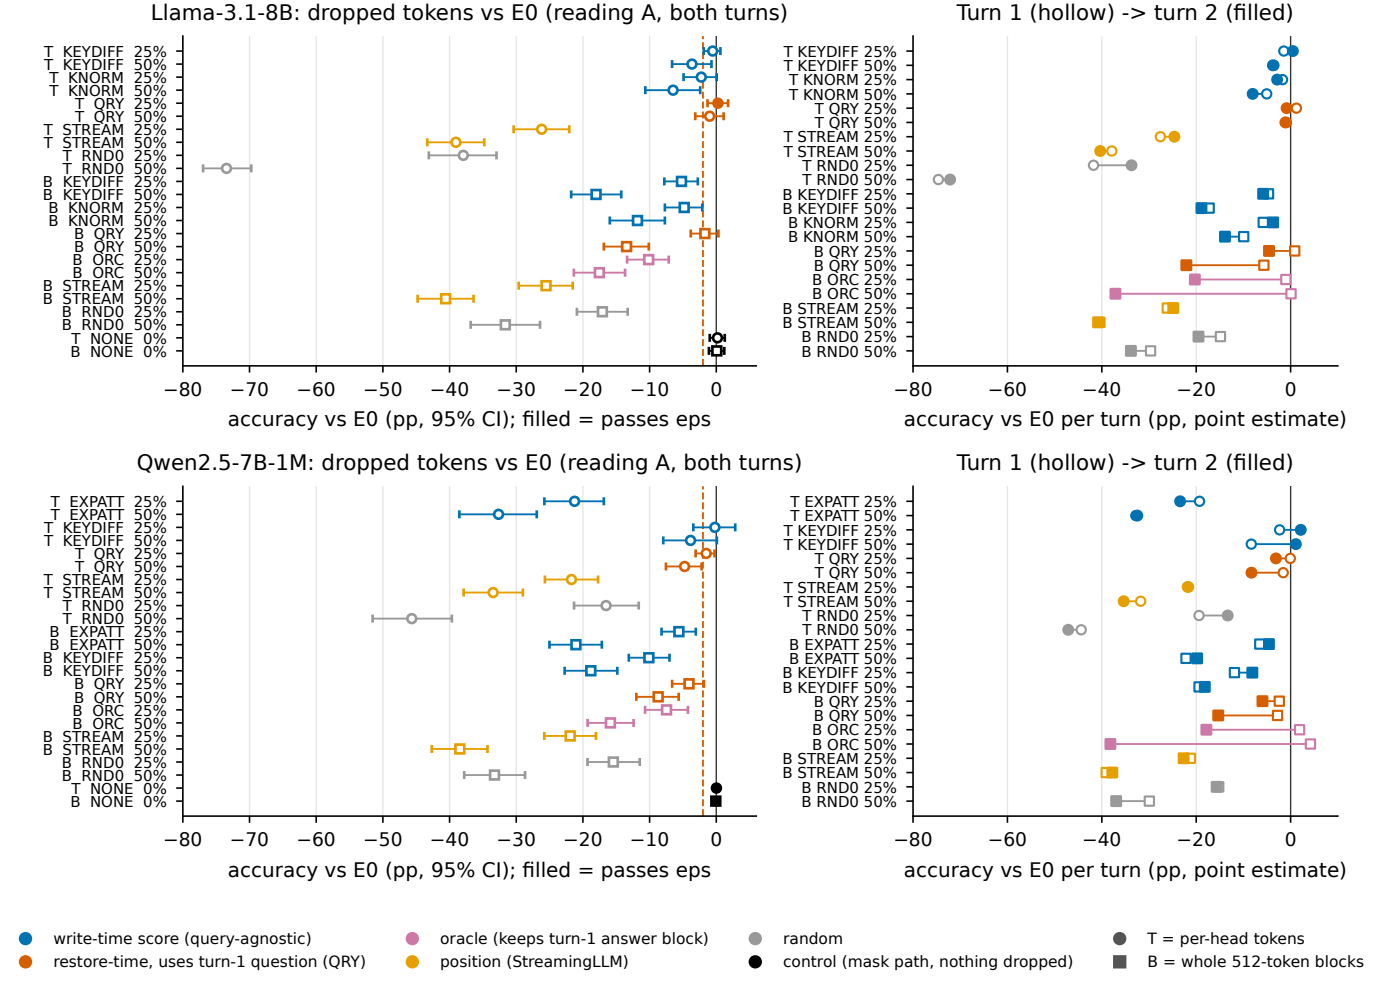
Llama-3.1-8B: dropped tokens vs E0 (reading A, both turns)
Turn 1 (hollow) -> turn 2 (filled)
T KEYDIFF 25%
T KEYDIFF 50%
T KNORM 25%
T KNORM 50%
T QRY 25%
T QRY 50%
T STREAM 25%
T STREAM 50%
T RND0 25%
T RND0 50%
B KEYDIFF 25%
B KEYDIFF 50%
B KNORM 25%
B KNORM 50%
B QRY 25%
B QRY 50%
B ORC 25%
B ORC 50%
B STREAM 25%
B STREAM 50%
B RND0 25%
B RND0 50%
T NONE 0%
B NONE 0%
−80 −70 −60 −50 −40 −30 −20 −10 0
accuracy vs E0 (pp, 95% CI); filled = passes eps
−80 −60 −40 −20 0
accuracy vs E0 per turn (pp, point estimate)
Qwen2.5-7B-1M: dropped tokens vs E0 (reading A, both turns)
Turn 1 (hollow) -> turn 2 (filled)
T EXPATT 25%
T EXPATT 50%
T KEYDIFF 25%
T KEYDIFF 50%
T QRY 25%
T QRY 50%
T STREAM 25%
T STREAM 50%
T RND0 25%
T RND0 50%
B EXPATT 25%
B EXPATT 50%
B KEYDIFF 25%
B KEYDIFF 50%
B QRY 25%
B QRY 50%
B ORC 25%
B ORC 50%
B STREAM 25%
B STREAM 50%
B RND0 25%
B RND0 50%
T NONE 0%
B NONE 0%
−80 −70 −60 −50 −40 −30 −20 −10 0
accuracy vs E0 (pp, 95% CI); filled = passes eps
−80 −60 −40 −20 0
accuracy vs E0 per turn (pp, point estimate)
write-time score (query-agnostic)
restore-time, uses turn-1 question (QRY)
oracle (keeps turn-1 answer block)
position (StreamingLLM)
random
control (mask path, nothing dropped)
T = per-head tokens
B = whole 512-token blocks