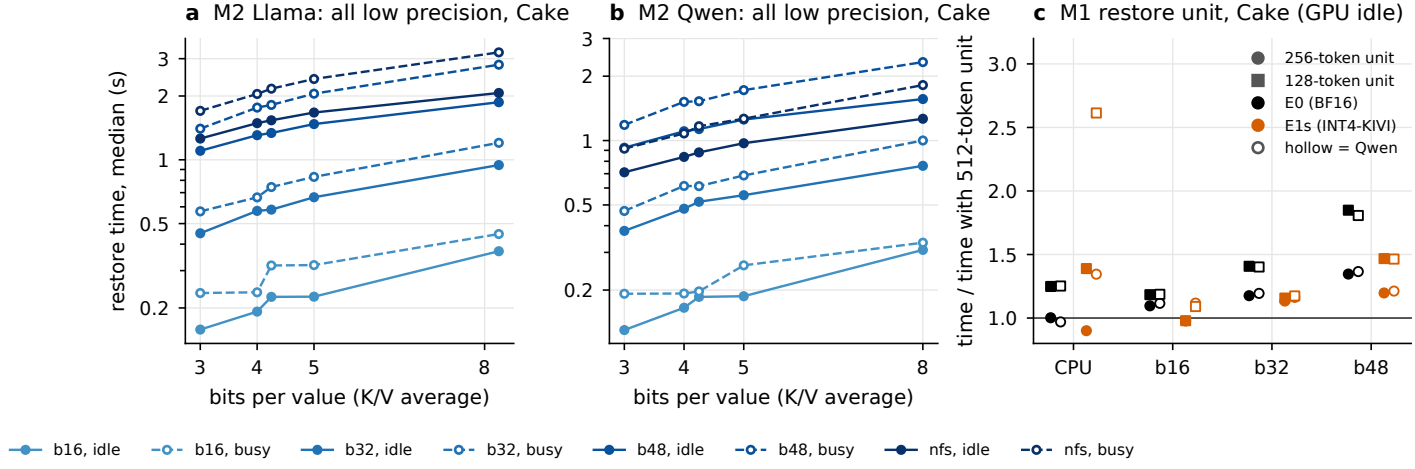
a M2 Llama: all low precision, Cake
b M2 Qwen: all low precision, Cake
c M1 restore unit, Cake (GPU idle)
restore time, median (s)
bits per value (K/V average)
time / time with 512-token unit
256-token unit
128-token unit
E0 (BF16)
E1s (INT4-KIVI)
hollow = Qwen
CPU
b16
b32
b48
b16, idle
b16, busy
b32, idle
b32, busy
b48, idle
b48, busy
nfs, idle
nfs, busy