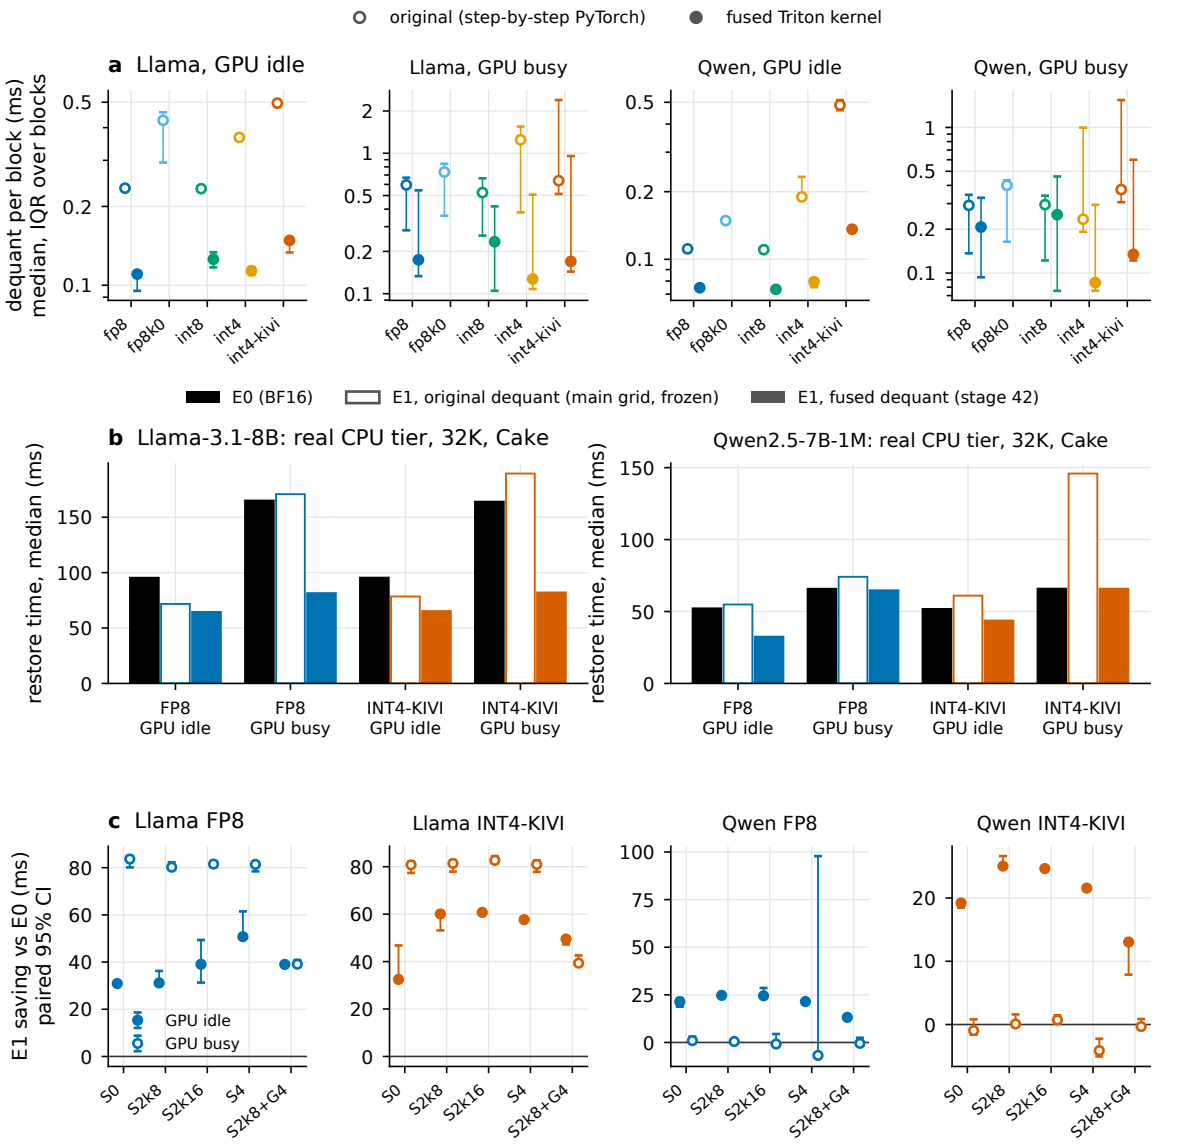


(c) 42c dequant-scheduling finals, real CPU tier, 32K, Cake: S0 per-block, S2k8/S2k16 = 8/16 blocks per kernel, S4 = all blocks once after loading, G4 = predicted meet (also speeds up E0, so its saving vs E0 looks smaller). All runs gpu_guard CLEAN.