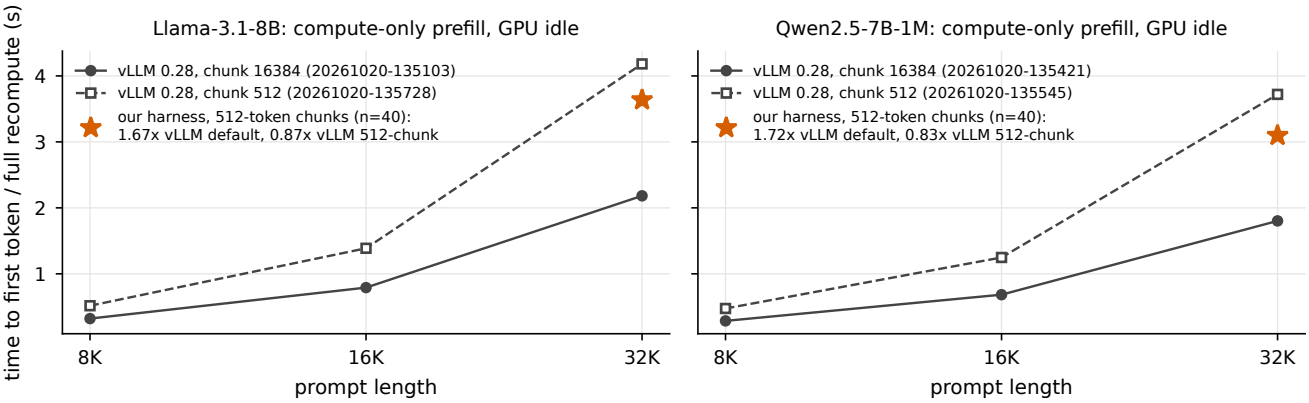
Llama-3.1-8B: compute-only prefill, GPU idle
vLLM 0.28, chunk 16384 (20261020-135103)
vLLM 0.28, chunk 512 (20261020-135728)
our harness, 512-token chunks (n=40):
1.67x vLLM default, 0.87x vLLM 512-chunk
time to first token / full recompute (s)
1
2
3
4
8K
16K
32K
prompt length
Qwen2.5-7B-1M: compute-only prefill, GPU idle
vLLM 0.28, chunk 16384 (20261020-135421)
vLLM 0.28, chunk 512 (20261020-135545)
our harness, 512-token chunks (n=40):
1.72x vLLM default, 0.83x vLLM 512-chunk
8K
16K
32K
prompt length